

```
from google.colab import drive
drive.mount('/content/drive')
```

Mounted at /content/drive

```
!pip install mlxtend
```

```
Looking in indexes: https://pypi.org/simple, https://us-python.pkg.dev/colab-wheels/
Requirement already satisfied: mlxtend in /usr/local/lib/python3.8/dist-packages (0.
Requirement already satisfied: matplotlib>=1.5.1 in /usr/local/lib/python3.8/dist-pa
Requirement already satisfied: setuptools in /usr/local/lib/python3.8/dist-packages
Requirement already satisfied: scikit-learn>=0.18 in /usr/local/lib/python3.8/dist-p
Requirement already satisfied: numpy>=1.10.4 in /usr/local/lib/python3.8/dist-packag
Requirement already satisfied: scipy>=0.17 in /usr/local/lib/python3.8/dist-packages
Requirement already satisfied: pandas>=0.17.1 in /usr/local/lib/python3.8/dist-packa
Requirement already satisfied: kiwisolver>=1.0.1 in /usr/local/lib/python3.8/dist-pa
Requirement already satisfied: cycloper>=0.10 in /usr/local/lib/python3.8/dist-package
Requirement already satisfied: python-dateutil>=2.1 in /usr/local/lib/python3.8/dist
Requirement already satisfied: pyparsing!=2.0.4,!=2.1.2,!=2.1.6,>=2.0.1 in /usr/loca
Requirement already satisfied: pytz>=2017.3 in /usr/local/lib/python3.8/dist-package
Requirement already satisfied: joblib>=0.11 in /usr/local/lib/python3.8/dist-package
Requirement already satisfied: threadpoolctl>=2.0.0 in /usr/local/lib/python3.8/dist
Requirement already satisfied: six>=1.5 in /usr/local/lib/python3.8/dist-packages (f
```

```
import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
```

```
import pickle
from os import path
```

```
from sklearn.preprocessing import MinMaxScaler
from sklearn.preprocessing import StandardScaler
from sklearn.preprocessing import LabelEncoder
```

```
from sklearn import preprocessing
import xgboost as xgb
from sklearn import svm
from sklearn.metrics import classification_report, accuracy_score, roc_auc_score, average_pre
import sklearn
from sklearn.model_selection import train_test_split
import matplotlib.pyplot as plt
import seaborn as sns
from yellowbrick.classifier import ClassificationReport
from sklearn.ensemble import AdaBoostClassifier, RandomForestClassifier
from sklearn.naive_bayes import GaussianNB
from sklearn.ensemble import VotingClassifier
import six
import sys
sys.modules['sklearn.externals.six'] = six
```

```

from mlxtend.classifier import StackingClassifier
from sklearn.neighbors import KNeighborsClassifier
from sklearn.linear_model import LogisticRegression
from sklearn.svm import SVC
from sklearn.model_selection import GridSearchCV
import joblib

sys.modules['sklearn.externals.joblib'] = joblib
from mlxtend.feature_selection import SequentialFeatureSelector as SFS
from sklearn.externals import joblib
import warnings
warnings.filterwarnings("ignore")

data = pd.read_csv('/content/drive/My Drive/NIDS_ML/UNSW_NB15_training-set.csv')
data.head()

```

| sport_ltm | ct_dst_src_ltm | is_ftp_login | ct_ftp_cmd | ct_flw_http_mthd | ct_src_ltm | ct_ |
|-----------|----------------|--------------|------------|------------------|------------|-----|
| 1         | 2              | 0            | 0          | 0                | 1          |     |
| 1         | 2              | 0            | 0          | 0                | 1          |     |
| 1         | 3              | 0            | 0          | 0                | 1          |     |
| 1         | 3              | 0            | 0          | 0                | 2          |     |
| 1         | 3              | 0            | 0          | 0                | 2          |     |

```
data.shape
```

```
(82332, 45)
```

```
data.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 82332 entries, 0 to 82331
Data columns (total 45 columns):
#   Column                Non-Null Count  Dtype
---  -
0   id                    82332 non-null  int64
1   dur                   82332 non-null  float64
2   proto                 82332 non-null  object
3   service               82332 non-null  object
4   state                 82332 non-null  object
5   spkts                 82332 non-null  int64
6   dpkts                 82332 non-null  int64
7   sbytes                82332 non-null  int64
8   dbytes                82332 non-null  int64
9   rate                  82332 non-null  float64
10  sttl                  82332 non-null  int64
11  dttl                  82332 non-null  int64

```

```

12  sload      82332 non-null float64
13  dload      82332 non-null float64
14  sloss      82332 non-null int64
15  dloss      82332 non-null int64
16  sinpkt     82332 non-null float64
17  dinpkt     82332 non-null float64
18  sjit       82332 non-null float64
19  djit       82332 non-null float64
20  swin       82332 non-null int64
21  stcpb      82332 non-null int64
22  dtcpb      82332 non-null int64
23  dwin       82332 non-null int64
24  tcprtt     82332 non-null float64
25  synack     82332 non-null float64
26  ackdat     82332 non-null float64
27  smean      82332 non-null int64
28  dmean      82332 non-null int64
29  trans_depth 82332 non-null int64
30  response_body_len 82332 non-null int64
31  ct_srv_src  82332 non-null int64
32  ct_state_ttl 82332 non-null int64
33  ct_dst_ltm  82332 non-null int64
34  ct_src_dport_ltm 82332 non-null int64
35  ct_dst_sport_ltm 82332 non-null int64
36  ct_dst_src_ltm 82332 non-null int64
37  is_ftp_login 82332 non-null int64
38  ct_ftp_cmd  82332 non-null int64
39  ct_flw_http_mthd 82332 non-null int64
40  ct_src_ltm  82332 non-null int64
41  ct_srv_dst  82332 non-null int64
42  is_sm_ips_ports 82332 non-null int64
43  attack_cat  82332 non-null object
44  label       82332 non-null int64
dtypes: float64(11), int64(30), object(4)
memory usage: 28.3+ MB

```

```
data.drop('service',axis='columns',inplace=True)
```

```
data.isnull().sum()
```

```

id      0
dur      0
proto    0
state    0
spkts    0
dpkts    0
sbytes   0
dbytes   0
rate     0
sttl     0
dttl     0
sload    0
dload    0
sloss    0
dloss    0
sinpkt   0
dinpkt   0
sjit     0

```

|                   |   |
|-------------------|---|
| djit              | 0 |
| swin              | 0 |
| stcpb             | 0 |
| dtcpb             | 0 |
| dwin              | 0 |
| tcprrt            | 0 |
| synack            | 0 |
| ackdat            | 0 |
| smean             | 0 |
| dmean             | 0 |
| trans_depth       | 0 |
| response_body_len | 0 |
| ct_srv_src        | 0 |
| ct_state_ttl      | 0 |
| ct_dst_ltm        | 0 |
| ct_src_dport_ltm  | 0 |
| ct_dst_sport_ltm  | 0 |
| ct_dst_src_ltm    | 0 |
| is_ftp_login      | 0 |
| ct_ftp_cmd        | 0 |
| ct_flw_http_mthd  | 0 |
| ct_src_ltm        | 0 |
| ct_srv_dst        | 0 |
| is_sm_ips_ports   | 0 |
| attack_cat        | 0 |
| label             | 0 |

dtype: int64

```
data['attack_cat'].value_counts()
```

|                |       |
|----------------|-------|
| Normal         | 37000 |
| Generic        | 18871 |
| Exploits       | 11132 |
| Fuzzers        | 6062  |
| DoS            | 4089  |
| Reconnaissance | 3496  |
| Analysis       | 677   |
| Backdoor       | 583   |
| Shellcode      | 378   |
| Worms          | 44    |

Name: attack\_cat, dtype: int64

```
data['state'].value_counts()
```

|     |       |
|-----|-------|
| FIN | 39339 |
| INT | 34163 |
| CON | 6982  |
| REQ | 1842  |
| ACC | 4     |
| RST | 1     |
| CLO | 1     |

Name: state, dtype: int64

```
data
```

|       | id    | dur      | proto | state | spkts | dpkts | sbytes | dbytes | rate          | st  |
|-------|-------|----------|-------|-------|-------|-------|--------|--------|---------------|-----|
| 0     | 1     | 0.000011 | udp   | INT   | 2     | 0     | 496    | 0      | 90909.090200  | 2   |
| 1     | 2     | 0.000008 | udp   | INT   | 2     | 0     | 1762   | 0      | 125000.000300 | 2   |
| 2     | 3     | 0.000005 | udp   | INT   | 2     | 0     | 1068   | 0      | 200000.005100 | 2   |
| 3     | 4     | 0.000006 | udp   | INT   | 2     | 0     | 900    | 0      | 166666.660800 | 2   |
| 4     | 5     | 0.000010 | udp   | INT   | 2     | 0     | 2126   | 0      | 100000.002500 | 2   |
| ...   | ...   | ...      | ...   | ...   | ...   | ...   | ...    | ...    | ...           | ... |
| 82327 | 82328 | 0.000005 | udp   | INT   | 2     | 0     | 104    | 0      | 200000.005100 | 2   |
| 82328 | 82329 | 1.106101 | tcp   | FIN   | 20    | 8     | 18062  | 354    | 24.410067     | 2   |
| 82329 | 82330 | 0.000000 | arp   | INT   | 1     | 0     | 46     | 0      | 0.000000      |     |
| 82330 | 82331 | 0.000000 | arp   | INT   | 1     | 0     | 46     | 0      | 0.000000      |     |
| 82331 | 82332 | 0.000009 | udp   | INT   | 2     | 0     | 104    | 0      | 111111.107200 | 2   |

82332 rows × 44 columns



```
features = pd.read_csv('/content/drive/MyDrive/NIDS_ML/NUSW-NB15_features.csv', encoding='c
```

```
features.head()
```

| No. | Name | Type   | Description |                         |
|-----|------|--------|-------------|-------------------------|
| 0   | 1    | srcip  | nominal     | Source IP address       |
| 1   | 2    | sport  | integer     | Source port number      |
| 2   | 3    | dstip  | nominal     | Destination IP address  |
| 3   | 4    | dsport | integer     | Destination port number |
| 4   | 5    | proto  | nominal     | Transaction protocol    |

```
features['Type '] = features['Type '].str.lower()
```

```
# selecting column names of all data types
nominal_names = features['Name'][features['Type ']=='nominal']
integer_names = features['Name'][features['Type ']=='integer']
binary_names = features['Name'][features['Type ']=='binary']
float_names = features['Name'][features['Type ']=='float']
```

```
cols = data.columns
nominal_names = cols.intersection(nominal_names)
integer_names = cols.intersection(integer_names)
binary_names = cols.intersection(binary_names)
float_names = cols.intersection(float_names)
```

```
for c in integer_names:
    pd.to_numeric(data[c])
```

```
for c in binary_names:
    pd.to_numeric(data[c])
```

```
for c in float_names:
    pd.to_numeric(data[c])
```

```
data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 82332 entries, 0 to 82331
Data columns (total 44 columns):
#   Column                Non-Null Count  Dtype
---  -
0   id                    82332 non-null  int64
1   dur                   82332 non-null  float64
2   proto                 82332 non-null  object
3   state                 82332 non-null  object
4   spkts                 82332 non-null  int64
5   dpkts                 82332 non-null  int64
6   sbytes                82332 non-null  int64
7   dbytes                82332 non-null  int64
8   rate                  82332 non-null  float64
9   sttl                  82332 non-null  int64
10  dttl                  82332 non-null  int64
11  sload                 82332 non-null  float64
12  dload                 82332 non-null  float64
13  sloss                 82332 non-null  int64
14  dloss                 82332 non-null  int64
15  sinpkt                82332 non-null  float64
16  dinpkt                82332 non-null  float64
17  sjit                  82332 non-null  float64
18  djit                  82332 non-null  float64
19  swin                  82332 non-null  int64
20  stcpb                 82332 non-null  int64
21  dtcpb                 82332 non-null  int64
22  dwin                  82332 non-null  int64
23  tcprtt                82332 non-null  float64
24  synack                 82332 non-null  float64
25  ackdat                 82332 non-null  float64
26  smean                 82332 non-null  int64
27  dmean                 82332 non-null  int64
28  trans_depth           82332 non-null  int64
29  response_body_len     82332 non-null  int64
30  ct_srv_src            82332 non-null  int64
31  ct_state_ttl          82332 non-null  int64
32  ct_dst_ltm            82332 non-null  int64
33  ct_src_dport_ltm     82332 non-null  int64
34  ct_dst_sport_ltm     82332 non-null  int64
35  ct_dst_src_ltm       82332 non-null  int64
36  is_ftp_login          82332 non-null  int64
37  ct_ftp_cmd            82332 non-null  int64
```

```

38 ct_flw_http_mthd 82332 non-null int64
39 ct_src_ltm       82332 non-null int64
40 ct_srv_dst       82332 non-null int64
41 is_sm_ips_ports  82332 non-null int64
42 attack_cat       82332 non-null object
43 label            82332 non-null int64
dtypes: float64(11), int64(30), object(3)
memory usage: 27.6+ MB

```

data

|              | id    | dur      | proto | state | spkts | dpkts | sbytes | dbytes | rate          | st  |
|--------------|-------|----------|-------|-------|-------|-------|--------|--------|---------------|-----|
| <b>0</b>     | 1     | 0.000011 | udp   | INT   | 2     | 0     | 496    | 0      | 90909.090200  | 2   |
| <b>1</b>     | 2     | 0.000008 | udp   | INT   | 2     | 0     | 1762   | 0      | 125000.000300 | 2   |
| <b>2</b>     | 3     | 0.000005 | udp   | INT   | 2     | 0     | 1068   | 0      | 200000.005100 | 2   |
| <b>3</b>     | 4     | 0.000006 | udp   | INT   | 2     | 0     | 900    | 0      | 166666.660800 | 2   |
| <b>4</b>     | 5     | 0.000010 | udp   | INT   | 2     | 0     | 2126   | 0      | 100000.002500 | 2   |
| ...          | ...   | ...      | ...   | ...   | ...   | ...   | ...    | ...    | ...           | ... |
| <b>82327</b> | 82328 | 0.000005 | udp   | INT   | 2     | 0     | 104    | 0      | 200000.005100 | 2   |
| <b>82328</b> | 82329 | 1.106101 | tcp   | FIN   | 20    | 8     | 18062  | 354    | 24.410067     | 2   |
| <b>82329</b> | 82330 | 0.000000 | arp   | INT   | 1     | 0     | 46     | 0      | 0.000000      |     |
| <b>82330</b> | 82331 | 0.000000 | arp   | INT   | 1     | 0     | 46     | 0      | 0.000000      |     |
| <b>82331</b> | 82332 | 0.000009 | udp   | INT   | 2     | 0     | 104    | 0      | 111111.107200 | 2   |

82332 rows × 44 columns

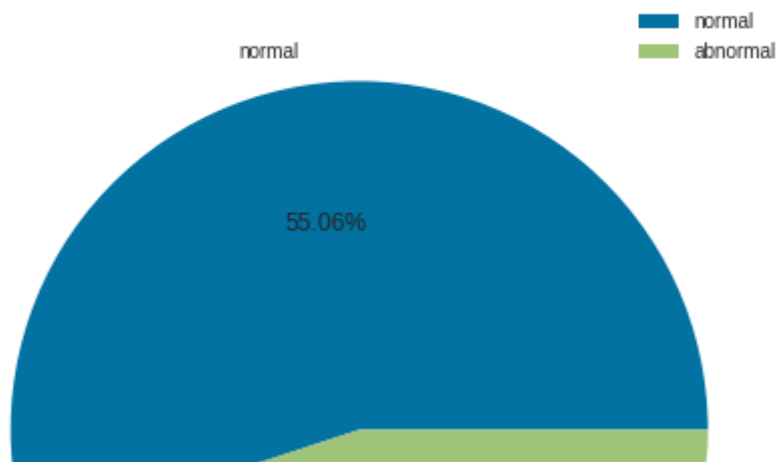


```

plt.figure(figsize=(8,8))
plt.pie(data.label.value_counts(),labels=['normal','abnormal'],autopct='%0.2f%%')
plt.title("Pie chart distribution of normal and abnormal labels",fontsize=16)
plt.legend()
plt.show()


```

Pie chart distribution of normal and abnormal labels



```
plt.figure(figsize=(8,8))
plt.pie(data.attack_cat.value_counts(),labels=data.attack_cat.unique(),autopct='%0.2f%%')
plt.title('Pie chart distribution of multi-class labels')
plt.legend(loc='best')
plt.show()
```


Pie chart distribution of multi-class labels



```
sns.countplot(data['label'])
```

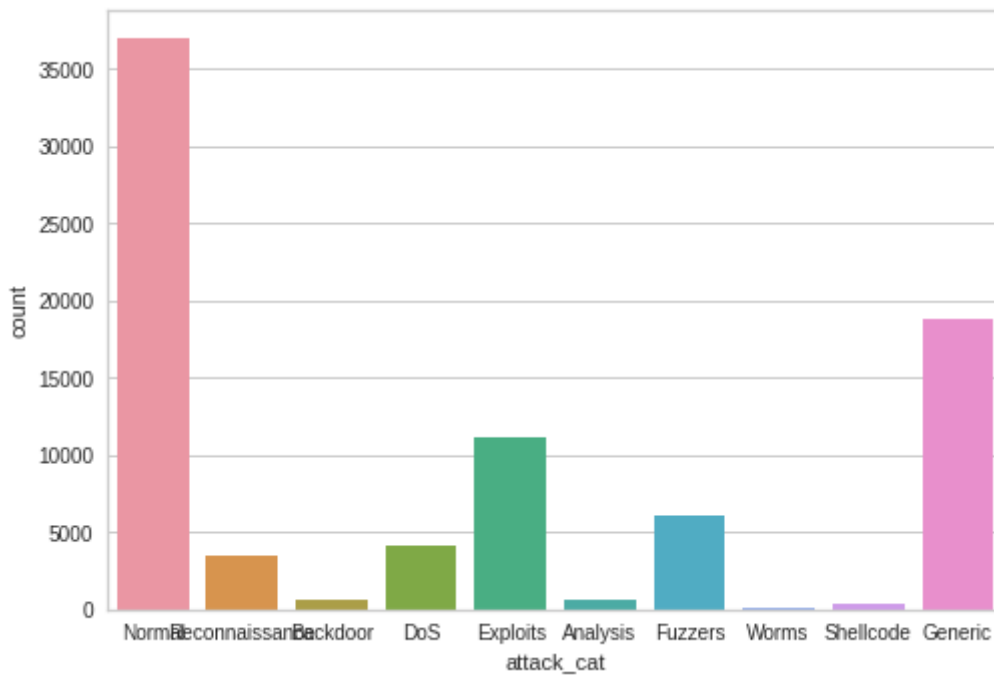


```
<matplotlib.axes._subplots.AxesSubplot at 0x7f5b2e6e99a0>
```



```
sns.countplot(data['attack_cat'])
```

```
<matplotlib.axes._subplots.AxesSubplot at 0x7f5b2e7494f0>
```




```
num_col = data.select_dtypes(include='number').columns
```

```
# selecting categorical data attributes
cat_col = data.columns.difference(num_col)
cat_col = cat_col[1:]
cat_col
```


```
Index(['proto', 'state'], dtype='object')
```

```
data_cat = data[cat_col].copy()
data_cat.head()
```

|   | proto | state |  |
|---|-------|-------|---|
| 0 | udp   | INT   |   |
| 1 | udp   | INT   |   |

```
from sklearn import preprocessing
le = preprocessing.LabelEncoder()
data_cat['proto'] = le.fit_transform(data_cat['proto'])
data_cat['state'] = le.fit_transform(data_cat['state'])
```

data\_cat

|       | proto | state |  |
|-------|-------|-------|---|
| 0     | 117   | 4     |   |
| 1     | 117   | 4     |   |
| 2     | 117   | 4     |   |
| 3     | 117   | 4     |   |
| 4     | 117   | 4     |   |
| ...   | ...   | ...   |   |
| 82327 | 117   | 4     |   |
| 82328 | 111   | 3     |   |
| 82329 | 6     | 4     |   |
| 82330 | 6     | 4     |   |
| 82331 | 117   | 4     |   |


82332 rows × 2 columns

```
data.drop(columns=cat_col,inplace=True)
```

```
data = pd.concat([data, data_cat],axis=1)
```

```
# selecting numeric attributes columns from data
num_col = list(data.select_dtypes(include='number').columns)
num_col.remove('id')
num_col.remove('label')
print(num_col)
```

```
, 'dbytes', 'rate', 'sttl', 'dttl', 'sload', 'dload', 'sloss', 'dloss', 'sinpkt', 'di
```



```
# using minmax scaler for normalizing data
minmax_scale = MinMaxScaler(feature_range=(0, 1))
def normalization(df,col):
```

```


for i in col:
    arr = df[i]
    arr = np.array(arr)
    df[i] = minmax_scale.fit_transform(arr.reshape(len(arr),1))
return df

```

```
data.head()
```

|   | id | dur      | spkts | dpkts | sbytes | dbytes | rate        | sttl | dttl | sload       |
|---|----|----------|-------|-------|--------|--------|-------------|------|------|-------------|
| 0 | 1  | 0.000011 | 2     | 0     | 496    | 0      | 90909.0902  | 254  | 0    | 180363632.0 |
| 1 | 2  | 0.000008 | 2     | 0     | 1762   | 0      | 125000.0003 | 254  | 0    | 881000000.0 |
| 2 | 3  | 0.000005 | 2     | 0     | 1068   | 0      | 200000.0051 | 254  | 0    | 854400000.0 |
| 3 | 4  | 0.000006 | 2     | 0     | 900    | 0      | 166666.6608 | 254  | 0    | 600000000.0 |
| 4 | 5  | 0.000010 | 2     | 0     | 2126   | 0      | 100000.0025 | 254  | 0    | 850400000.0 |

5 rows × 44 columns



```
data = normalization(data.copy(),num_col)
```

```
data.head()
```

| load | ... | is_ftp_login | ct_ftp_cmd | ct_flw_http_mthd | ct_src_ltm | ct_srv_dst | is_sm_i |
|------|-----|--------------|------------|------------------|------------|------------|---------|
| 1238 | ... | 0.0          | 0.0        | 0.0              | 0.000000   | 0.016393   |         |
| 7236 | ... | 0.0          | 0.0        | 0.0              | 0.000000   | 0.016393   |         |
| 2187 | ... | 0.0          | 0.0        | 0.0              | 0.000000   | 0.032787   |         |
| 3895 | ... | 0.0          | 0.0        | 0.0              | 0.016949   | 0.032787   |         |
| 1427 | ... | 0.0          | 0.0        | 0.0              | 0.016949   | 0.032787   |         |



## Binary Labels

```
bin_label = pd.DataFrame(data.label.map(lambda x:'normal' if x==0 else 'abnormal'))

# creating a dataframe with binary labels (normal,abnormal)
bin_data = data.copy()
bin_data['label'] = bin_label

# label encoding (0,1) binary labels
le1 = preprocessing.LabelEncoder()
enc_label = bin_label.apply(le1.fit_transform)
bin_data['label'] = enc_label

le1.classes_

array(['abnormal', 'normal'], dtype=object)

np.save("le1_classes.npy",le1.classes_,allow_pickle=True)
```

## Multi-class Labels


```
multi_data = data.copy()
multi_label = pd.DataFrame(multi_data.attack_cat)
multi_data = pd.get_dummies(multi_data,columns=['attack_cat'])
le2 = preprocessing.LabelEncoder()
enc_label = multi_label.apply(le2.fit_transform)
multi_data['label'] = enc_label
le2.classes_

array(['Analysis', 'Backdoor', 'DoS', 'Exploits', 'Fuzzers', 'Generic',
      'Normal', 'Reconnaissance', 'Shellcode', 'Worms'], dtype=object)


np.save("le2_classes.npy",le2.classes_,allow_pickle=True)

num_col.append('label')
```

```
# Correlation Matrix for Binary Labels
plt.figure(figsize=(20,8))
corr_bin = bin_data[num_col].corr()
sns.heatmap(corr_bin,vmax=1.0,annot=False)
plt.title('Correlation Matrix for Binary Labels',fontsize=16)
plt.show()
```



```
num_col = list(multi_data.select_dtypes(include='number').columns)
# Correlation Matrix for Multi-class Labels
plt.figure(figsize=(20,8))
corr_multi = multi_data[num_col].corr()
sns.heatmap(corr_multi,vmax=1.0,annot=False)
plt.title('Correlation Matrix for Multi Labels',fontsize=16)
plt.show()
```



## Feature Selection

### Binary Labels

```
corr_ybin = abs(corr_bin['label'])
highest_corr_bin = corr_ybin[corr_ybin > 0.3]
highest_corr_bin.sort_values(ascending=True)
```

```
ct_state_ttl      0.318517
rate              0.328629
ct_src_dport_ltm  0.341513
dwin              0.369257
ct_dst_sport_ltm  0.393668
swin              0.414504
state             0.459040
sttl              0.504159
label             1.000000
Name: label, dtype: float64
```

```
bin_cols = highest_corr_bin.index
bin_cols
```

```
Index(['rate', 'sttl', 'swin', 'dwin', 'ct_state_ttl', 'ct_src_dport_ltm',
      'ct_dst_sport_ltm', 'state', 'label'],
      dtype='object')
```

```
bin_data = bin_data[bin_cols].copy()
bin_data
```

|              | rate     | sttl     | swin | dwin | ct_state_ttl | ct_src_dport_ltm | ct_dst_sport_l |
|--------------|----------|----------|------|------|--------------|------------------|----------------|
| <b>0</b>     | 0.090909 | 0.996078 | 0.0  | 0.0  | 0.333333     | 0.000000         |                |
| <b>1</b>     | 0.125000 | 0.996078 | 0.0  | 0.0  | 0.333333     | 0.000000         |                |
| <b>2</b>     | 0.200000 | 0.996078 | 0.0  | 0.0  | 0.333333     | 0.000000         |                |
| <b>3</b>     | 0.166667 | 0.996078 | 0.0  | 0.0  | 0.333333     | 0.017241         |                |
| <b>4</b>     | 0.100000 | 0.996078 | 0.0  | 0.0  | 0.333333     | 0.017241         |                |
| ...          | ...      | ...      | ...  | ...  | ...          | ...              | ...            |
| <b>82327</b> | 0.200000 | 0.996078 | 0.0  | 0.0  | 0.333333     | 0.000000         |                |
| <b>82328</b> | 0.000024 | 0.996078 | 1.0  | 1.0  | 0.166667     | 0.000000         |                |
| <b>82329</b> | 0.000000 | 0.000000 | 0.0  | 0.0  | 0.333333     | 0.000000         |                |
| <b>82330</b> | 0.000000 | 0.000000 | 0.0  | 0.0  | 0.333333     | 0.000000         |                |
| <b>82331</b> | 0.111111 | 0.996078 | 0.0  | 0.0  | 0.333333     | 0.000000         |                |

82332 rows × 9 columns



```
bin_data.to_csv('./bin_data.csv')
```

## Multi-class Labels

```
# finding the attributes which have more than 0.3 correlation with encoded attack label at
corr_ymulti = abs(corr_multi['label'])
highest_corr_multi = corr_ymulti[corr_ymulti > 0.2]
highest_corr_multi.sort_values(ascending=True)
```

```
state          0.214254
attack_cat_Backdoor  0.235245
id             0.274428
attack_cat_Reconnaissance  0.296008
attack_cat_Analysis  0.317254
attack_cat_DoS      0.477123
attack_cat_Exploits  0.549046
attack_cat_Normal    0.638825
label            1.000000
Name: label, dtype: float64
```

```
# selecting attributes found by using pearson correlation coefficient
multi_cols = highest_corr_multi.index
multi_cols
```

```
Index(['id', 'label', 'state', 'attack_cat_Analysis', 'attack_cat_Backdoor',
      'attack_cat_DoS', 'attack_cat_Exploits', 'attack_cat_Normal',
      'attack_cat_Reconnaissance'],
      dtype='object')
```

```
multi_data = multi_data[multi_cols].copy()
```

```
multi_data.to_csv('./multi_data.csv')
```

## BINARY CLASSIFICATION

```
X = bin_data.drop(columns=['label'],axis=1)
Y = bin_data['label']
```

```
X_train,X_test,y_train,y_test = train_test_split(X,Y,test_size=0.20, random_state=50)
```

```
#GaussianNB
```

```
def GNB(X_train,y_train,X_test,y_test):
    gnb_clf = GaussianNB()
    pred = gnb_clf.fit(X_train, y_train).predict(X_test)
    pred= gnb_clf.predict(X_test)
    print ("GaussianNB:Accuracy : ", accuracy_score(y_test,pred)*100)
```

```
#confusion Matrix
```


```
matrix =confusion_matrix(y_test, pred)
sns.heatmap(pd.DataFrame(matrix).annot=True, cmap="YlGnBu" ,fmt='g')
```

```
plt.tight_layout()
plt.title('Confusion matrix', y=1.1)
plt.ylabel('Actual label')
plt.xlabel('Predicted label')
plt.show()

#Classification Report
prediction=gnb_clf.predict(X_test)
print(classification_report(y_test, prediction))
visualizer = ClassificationReport(gnb_clf, support=True)
visualizer.fit(X_train, y_train)
visualizer.score(X_test, y_test)
g = visualizer.poof()
GNB(X_train,y_train,X_test,y_test)
```



GaussianNB:Accuracy : 75.55110220440882



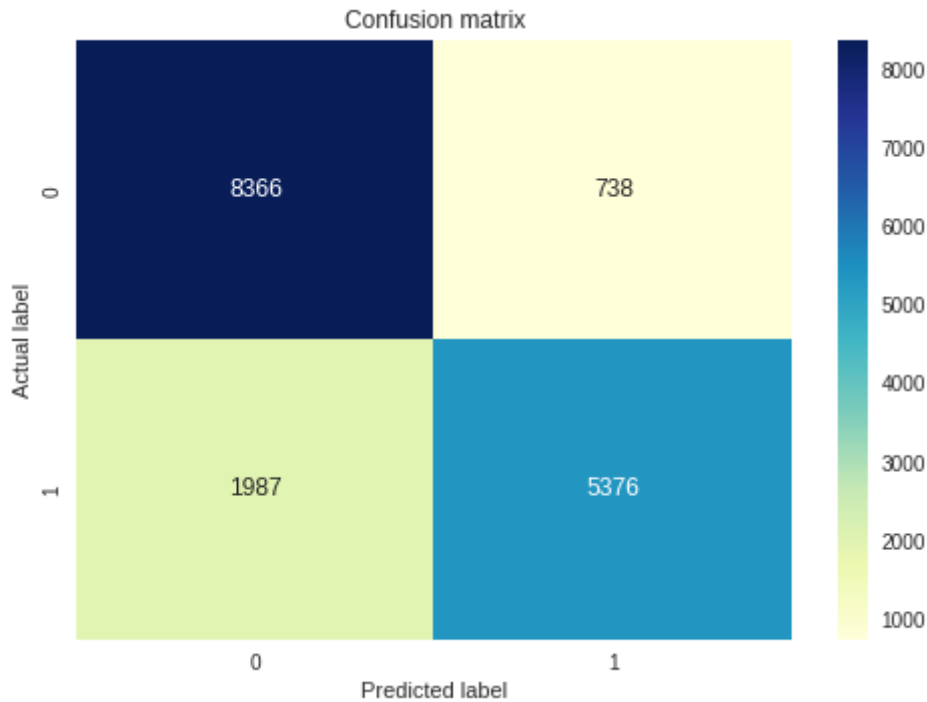
#KNeighbours

```
def KNN1(X_train,y_train,X_test,y_test):
    knn = KNeighborsClassifier(n_neighbors=2)
    knn.fit(X_train, y_train)
    pred = knn.predict(X_test)
    print ("KNN:Accuracy : ", accuracy_score(y_test,pred)*100)


    #confusion Matrix
    matrix =confusion_matrix(y_test, pred)
    sns.heatmap(pd.DataFrame(matrix), annot=True, cmap="YlGnBu" ,fmt='g')
    plt.title('Confusion matrix', y=1.1)
    plt.ylabel('Actual label')
    plt.xlabel('Predicted label')
    plt.show()

    #Classification Report
    prediction=knn.predict(X_test)
    print(classification_report(y_test, prediction))
    visualizer = ClassificationReport(knn, support=True)
    visualizer.fit(X_train, y_train)
    visualizer.score(X_test, y_test)
    g = visualizer.poof()
    KNN1(X_train,y_train,X_test,y_test)
```

KNN:Accuracy : 83.45175198882615



|              | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0            | 0.81      | 0.92   | 0.86     | 9104    |
| 1            | 0.88      | 0.73   | 0.80     | 7363    |
| accuracy     |           |        | 0.83     | 16467   |
| macro avg    | 0.84      | 0.82   | 0.83     | 16467   |
| weighted avg | 0.84      | 0.83   | 0.83     | 16467   |



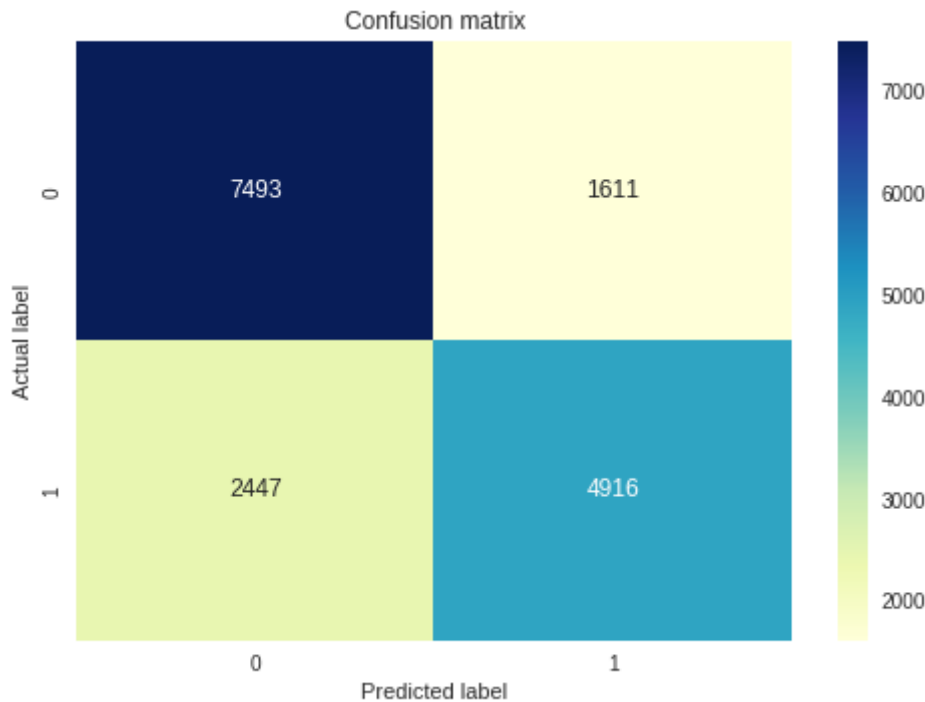
```
#LogisticRegression
from sklearn.linear_model import LogisticRegression
def logisticreg(X_train,y_train,X_test,y_test):
    lr = LogisticRegression()
    lr.fit(X_train,y_train)
    pred = lr.predict(X_test)
    print ("LogisticRegression:Accuracy : ", accuracy_score(y_test,pred)*100)

#confusion Matrix
matrix =confusion_matrix(y_test, pred)
sns.heatmap(pd.DataFrame(matrix), annot=True, cmap="YlGnBu" ,fmt='g')
plt.title('Confusion matrix', y=1.1)
plt.ylabel('Actual label')
plt.xlabel('Predicted label')
plt.show()

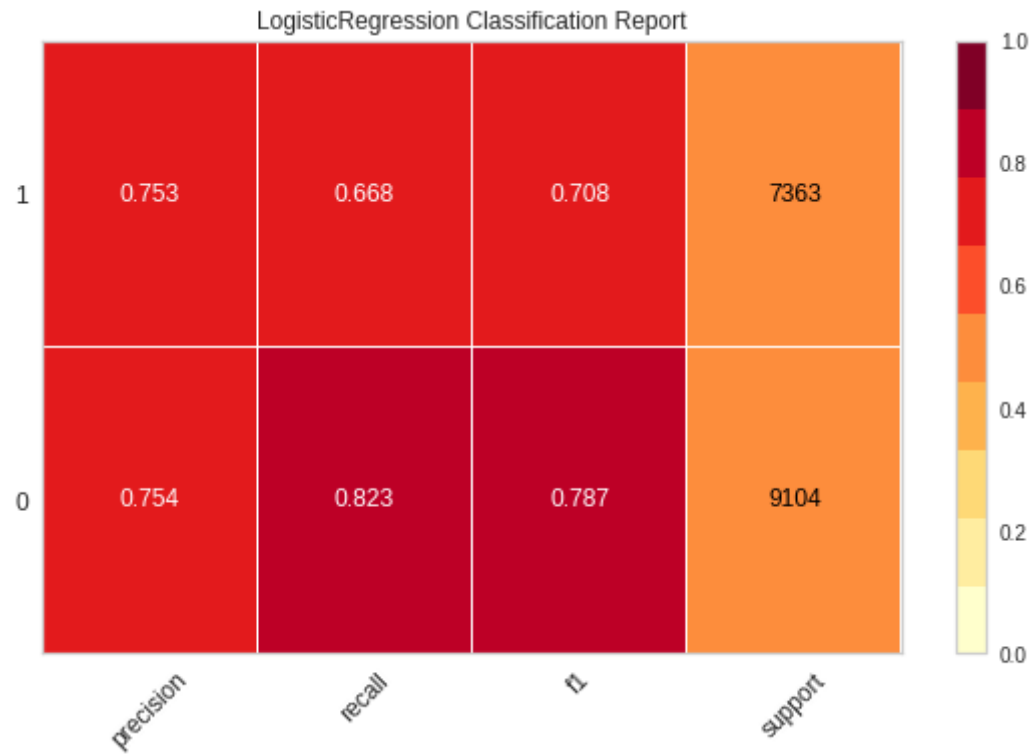
#Classification Report
prediction=lr.predict(X_test)
print(classification_report(y_test, prediction))
visualizer = ClassificationReport(lr, support=True)
visualizer.fit(X_train, y_train)
visualizer.score(X_test, y_test)
```

```
g = visualizer.poof()  
logisticreg(X_train,y_train,X_test,y_test)
```

LogisticRegression:Accuracy : 75.35677415436935



|              | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0            | 0.75      | 0.82   | 0.79     | 9104    |
| 1            | 0.75      | 0.67   | 0.71     | 7363    |
| accuracy     |           |        | 0.75     | 16467   |
| macro avg    | 0.75      | 0.75   | 0.75     | 16467   |
| weighted avg | 0.75      | 0.75   | 0.75     | 16467   |



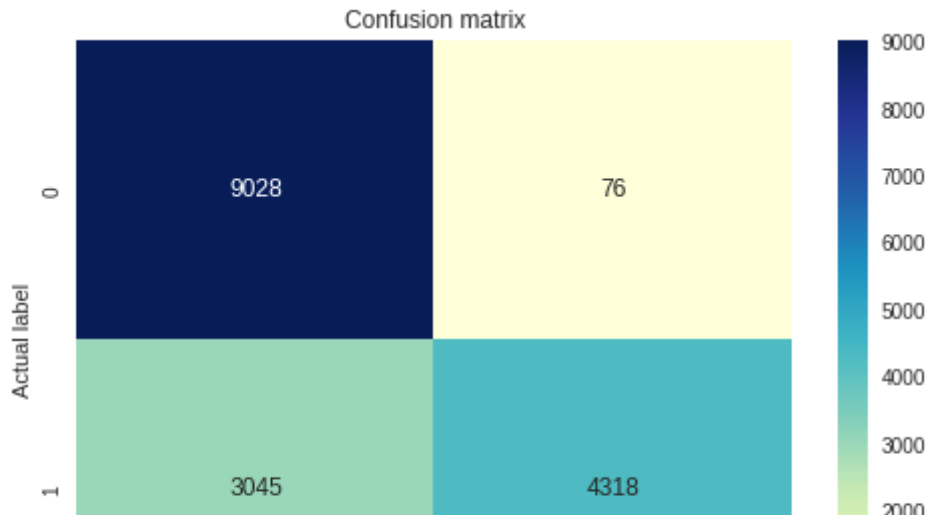
```
#Randon Forest  
def random_forest(X_train,y_train,X_test,y_test):  
    rf = RandomForestClassifier(max_depth=2, min_samples_split=2)  
    rf.fit(X_train,y_train)
```

```
rf.fit(X_train,y_train)
pred = rf.predict(X_test)
print ("Random Forest:Accuracy : ", accuracy_score(y_test,pred)*100)

#confusion Matrix
matrix =confusion_matrix(y_test, pred)
sns.heatmap(pd.DataFrame(matrix), annot=True, cmap="YlGnBu" ,fmt='g')
plt.title('Confusion matrix', y=1.1)
plt.ylabel('Actual label')
plt.xlabel('Predicted label')
plt.show()

#Classification Report
prediction=rf.predict(X_test)
print(classification_report(y_test, prediction))
visualizer = ClassificationReport(rf, support=True)
visualizer.fit(X_train, y_train)
visualizer.score(X_test, y_test)
g = visualizer.poof()
random_forest(X_train,y_train,X_test,y_test)
```

Random Forest:Accuracy : 81.04694236958765



#Stacking Classifier

```
xg = xgb.XGBClassifier(max_depth=5, learning_rate=0.01, n_estimators=100, gamma=0,min_child_weight=1)
rf = RandomForestClassifier(bootstrap=True,max_depth= 70,max_features= 'auto',min_samples_split=10,min_samples_leaf=10)
knn=KNeighborsClassifier()
```

```
def stacking(X_train,y_train,X_test,y_test):
    classifiers=[rf,knn]
    sc = StackingClassifier(classifiers,meta_classifier=xg)
    sc.fit(X_train,y_train)
    pred = sc.predict(X_test)
    print ("Stacking Classifier:Accuracy : ", accuracy_score(y_test,pred)*100)
```

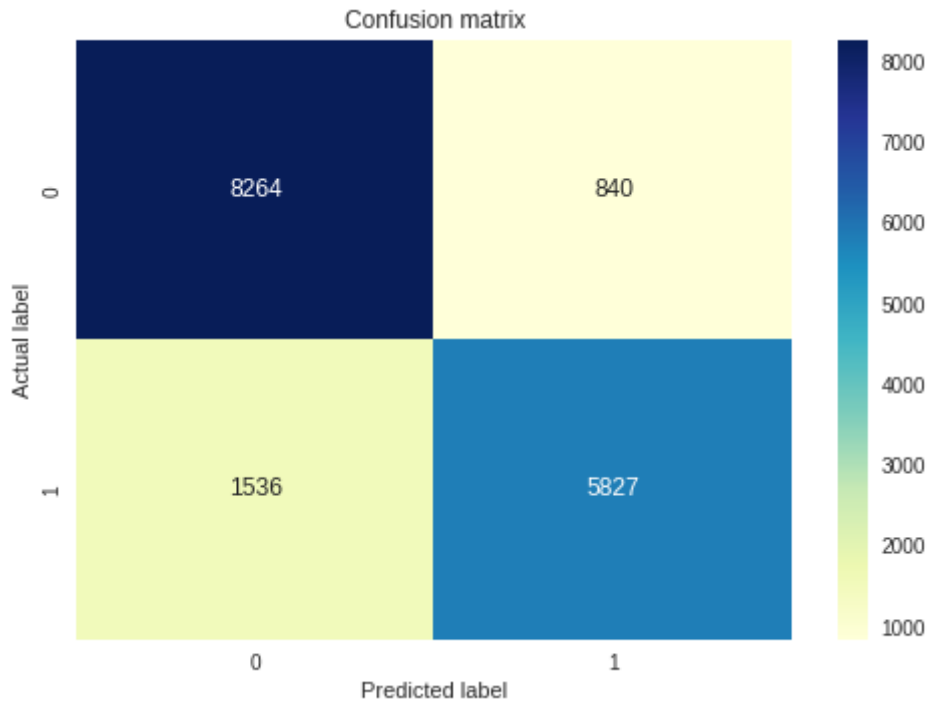
#confusion Matrix

```
matrix =confusion_matrix(y_test, pred)
sns.heatmap(pd.DataFrame(matrix), annot=True, cmap="YlGnBu" ,fmt='g')
plt.title('Confusion matrix', y=1.1)
plt.ylabel('Actual label')
plt.xlabel('Predicted label')
plt.show()
```


#Classification Report

```
prediction=sc.predict(X_test)
print(classification_report(y_test, prediction))
visualizer = ClassificationReport(sc, support=True)
visualizer.fit(X_train, y_train)
visualizer.score(X_test, y_test)
g = visualizer.poof()
stacking(X_train,y_train,X_test,y_test)
```

Stacking Classifier:Accuracy : 85.57114228456913



|              | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0            | 0.84      | 0.91   | 0.87     | 9104    |
| 1            | 0.87      | 0.79   | 0.83     | 7363    |
| accuracy     |           |        | 0.86     | 16467   |
| macro avg    | 0.86      | 0.85   | 0.85     | 16467   |
| weighted avg | 0.86      | 0.86   | 0.85     | 16467   |



## MULTI-CLASS CLASSIFICATION



```

X = multi_data.drop(columns=['label'],axis=1)
Y = multi_data['label']
X_train,X_test,y_train,y_test = train_test_split(X,Y,test_size=0.30, random_state=100)


#GaussianNB
def GNB(X_train,y_train,X_test,y_test):
    gnb_clf = GaussianNB()
    pred = gnb_clf.fit(X_train, y_train).predict(X_test)
    pred= gnb_clf.predict(X_test)

    print ("GaussianNB:Accuracy : ", accuracy_score(y_test,pred)*100)
  
```

```
#confusion Matrix
matrix =confusion_matrix(y_test, pred)
sns.heatmap(pd.DataFrame(matrix), annot=True, cmap="YlGnBu" ,fmt='g')
plt.tight_layout()
plt.title('Confusion matrix', y=1.1)
plt.ylabel('Actual label')
plt.xlabel('Predicted label')
plt.show()

#Classification Report
prediction=gnb_clf.predict(X_test)
print(classification_report(y_test, prediction))
visualizer = ClassificationReport(gnb_clf, support=True)
visualizer.fit(X_train, y_train)
visualizer.score(X_test, y_test)
g = visualizer.poof()
GNB(X_train,y_train,X_test,y_test)
```

GaussianNB:Accuracy : 71.33198380566802



#KNeighbours

```
def KNN1(X_train,y_train,X_test,y_test):
```

```
    knn = KNeighborsClassifier()
```

```
    knn.fit(X_train, y_train)
```

```
    pred = knn.predict(X_test)
```

```
    print ("KNN:Accuracy : ", accuracy_score(y_test,pred)*100)
```

```
    #confusion Matrix
```

```
    matrix =confusion_matrix(y_test, pred)
```

```
    sns.heatmap(pd.DataFrame(matrix), annot=True, cmap="YlGnBu" ,fmt='g')
```

```
    plt.title('Confusion matrix', y=1.1)
```

```
    plt.ylabel('Actual label')
```

```
    plt.xlabel('Predicted label')
```

```
    plt.show()
```

```
    #Classification Report
```

```
    prediction=knn.predict(X_test)
```

```
    print(classification_report(y_test, prediction))
```

```
    visualizer = ClassificationReport(knn, support=True)
```

```
    visualizer.fit(X_train, y_train)
```


```
    visualizer.score(X_test, y_test)
```

```
    g = visualizer.poof()
```

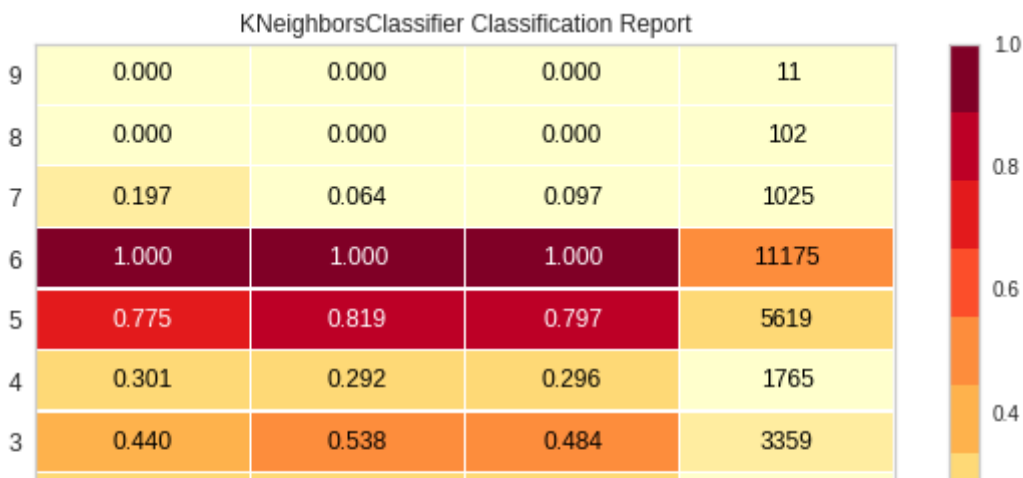
```
    KNN1(X_train,y_train,X_test,y_test)
```



KNN:Accuracy : 75.10121457489879



|              | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0            | 0.04      | 0.03   | 0.03     | 200     |
| 1            | 0.02      | 0.01   | 0.01     | 177     |
| 2            | 0.31      | 0.30   | 0.30     | 1267    |
| 3            | 0.44      | 0.54   | 0.48     | 3359    |
| 4            | 0.30      | 0.29   | 0.30     | 1765    |
| 5            | 0.78      | 0.82   | 0.80     | 5619    |
| 6            | 1.00      | 1.00   | 1.00     | 11175   |
| 7            | 0.20      | 0.06   | 0.10     | 1025    |
| 8            | 0.00      | 0.00   | 0.00     | 102     |
| 9            | 0.00      | 0.00   | 0.00     | 11      |
| accuracy     |           |        | 0.75     | 24700   |
| macro avg    | 0.31      | 0.30   | 0.30     | 24700   |
| weighted avg | 0.73      | 0.75   | 0.74     | 24700   |




```
#LogisticRegression
from sklearn.linear_model import LogisticRegression
def logisticreg(X_train,y_train,X_test,y_test):
    lr = LogisticRegression()
    lr.fit(X_train,y_train)
    pred = lr.predict(X_test)
```

```
print ("LogisticRegression:Accuracy : ", accuracy_score(y_test,pred)*100)

#confusion Matrix
matrix =confusion_matrix(y_test, pred)
sns.heatmap(pd.DataFrame(matrix), annot=True, cmap="YlGnBu" ,fmt='g')
plt.title('Confusion matrix', y=1.1)
plt.ylabel('Actual label')
plt.xlabel('Predicted label')
plt.show()

#Classification Report
prediction=lr.predict(X_test)
print(classification_report(y_test, prediction))
visualizer = ClassificationReport(lr, support=True)
visualizer.fit(X_train, y_train)
visualizer.score(X_test, y_test)
g = visualizer.poof()
logisticreg(X_train,y_train,X_test,y_test)
```

LogisticRegression:Accuracy : 45.24291497975709




#Random Forest

```
def random_forest(X_train,y_train,X_test,y_test):
    rf = RandomForestClassifier(max_depth=2, min_samples_split=2,)
    rf.fit(X_train,y_train)
    pred = rf.predict(X_test)
    print ("Random Forest:Accuracy : ", accuracy_score(y_test,pred)*100)

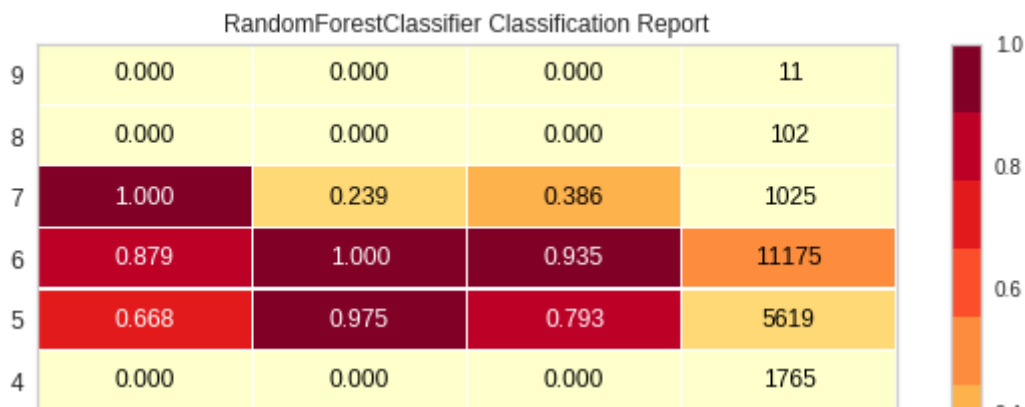
    #confusion Matrix
    matrix =confusion_matrix(y_test, pred)
    sns.heatmap(pd.DataFrame(matrix), annot=True, cmap="YlGnBu" ,fmt='g')
    plt.title('Confusion matrix', y=1.1)
    plt.ylabel('Actual label')
    plt.xlabel('Predicted label')
    plt.show()

    #Classification Report
    prediction=rf.predict(X_test)
    print(classification_report(y_test, prediction))
    visualizer = ClassificationReport(rf, support=True)
    visualizer.fit(X_train, y_train)
    visualizer.score(X_test, y_test)
    g = visualizer.poof()
    random_forest(X_train,y_train,X_test,y_test)
```

Random Forest:Accuracy : 82.75303643724696



|              | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0            | 0.00      | 0.00   | 0.00     | 200     |
| 1            | 0.00      | 0.00   | 0.00     | 177     |
| 2            | 1.00      | 0.14   | 0.25     | 1267    |
| 3            | 1.00      | 1.00   | 1.00     | 3359    |
| 4            | 0.00      | 0.00   | 0.00     | 1765    |
| 5            | 0.67      | 0.98   | 0.79     | 5619    |
| 6            | 0.88      | 1.00   | 0.94     | 11175   |
| 7            | 1.00      | 0.24   | 0.39     | 1025    |
| 8            | 0.00      | 0.00   | 0.00     | 102     |
| 9            | 0.00      | 0.00   | 0.00     | 11      |
| accuracy     |           |        | 0.83     | 24700   |
| macro avg    | 0.45      | 0.34   | 0.34     | 24700   |
| weighted avg | 0.78      | 0.83   | 0.77     | 24700   |



#Stacking Classifier

```

xg = xgb.XGBClassifier(max_depth=5, learning_rate=0.01, n_estimators=100, gamma=0,min_chi
rf = RandomForestClassifier(bootstrap=True,max_depth= 70,max_features= 'auto',min_samples
knn=KNeighborsClassifier())

```

```

def stacking(X_train,y_train,X_test,y_test):
    classifiers=[rf,knn]
    sc = StackingClassifier(classifiers,meta_classifier=xg)


```

```
sc.fit(X_train,y_train)
pred = sc.predict(X_test)
print ("Stacking Classifier:Accuracy : ", accuracy_score(y_test,pred)*100)

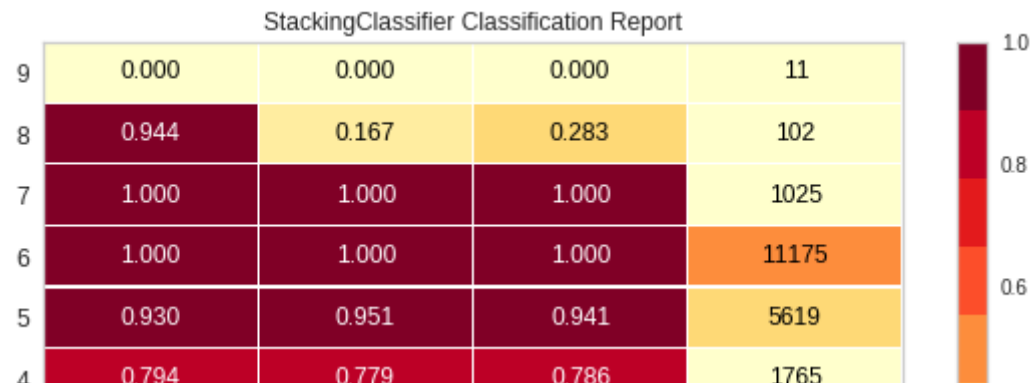
#confusion Matrix
matrix =confusion_matrix(y_test, pred)
sns.heatmap(pd.DataFrame(matrix), annot=True, cmap="YlGnBu" ,fmt='g')
plt.title('Confusion matrix', y=1.1)
plt.ylabel('Actual label')
plt.xlabel('Predicted label')
plt.show()

#Classification Report
prediction=sc.predict(X_test)
print(classification_report(y_test, prediction))
visualizer = ClassificationReport(sc, support=True)
visualizer.fit(X_train, y_train)
visualizer.score(X_test, y_test)
g = visualizer.poof()
stacking(X_train,y_train,X_test,y_test)
```

Stacking Classifier:Accuracy : 96.92307692307692



|              | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0            | 1.00      | 1.00   | 1.00     | 200     |
| 1            | 1.00      | 1.00   | 1.00     | 177     |
| 2            | 1.00      | 1.00   | 1.00     | 1267    |
| 3            | 1.00      | 1.00   | 1.00     | 3359    |
| 4            | 0.79      | 0.78   | 0.79     | 1765    |
| 5            | 0.93      | 0.95   | 0.94     | 5619    |
| 6            | 1.00      | 1.00   | 1.00     | 11175   |
| 7            | 1.00      | 1.00   | 1.00     | 1025    |
| 8            | 0.94      | 0.17   | 0.28     | 102     |
| 9            | 0.00      | 0.00   | 0.00     | 11      |
| accuracy     |           |        | 0.97     | 24700   |
| macro avg    | 0.87      | 0.79   | 0.80     | 24700   |
| weighted avg | 0.97      | 0.97   | 0.97     | 24700   |



✓ 35s

completed at 12:45 PM

● ✕